Whose Beauty?

Results-driven supervision update

Geographic and cultural conditioning effects in multimodal AI descriptions

Gemini-only pilot | 10 images | 150 responses | NLP + statistics

Focus today: results, statistical evidence, reproducible pipeline, and publication potential.

PROGRESS EVIDENCE

What has been delivered against the supervision brief

Supervisor requirement	Notebook evidence / output
Freeze design	One multimodal model; one prompt template; no-country control; four country labels
Proper pilot	10 images × 5 conditions × 3 repeats = 150 model responses
Not only manual reading	NLP variables extracted; summary tables; figures; statistical tests
Scientific measures	Semantic shift, sentiment, adjective categories, unsupported inference
Hypothesis link	Each measure defined as a testable indicator of geographic-conditioning effects
Produce results	Figures, tables, representative examples, Friedman and Wilcoxon results
Reproducibility	Raw responses, processed data, figures, tables, examples and protocol saved

Frozen experimental design: controlled repeated-measures pilot

1	10	5	3	150
multimodal model Gemini 3.5 Flash	images analysed	conditions per image	repeats per condition	total responses

Component	Final configuration
Control	No-country image description prompt
Country labels	Nigeria, India, Saudi Arabia, United Kingdom
Unit of analysis	Generated response, grouped by image × condition × repeat
Comparison logic	Country-conditioned output compared against the corresponding no-country control
Design rationale	Within-image comparison reduces image-to-image confounding

Reproducible analysis pipeline now exists

1. Input images **10 pilot images**	2. Gemini generation 5 conditions × 3 repeats	3. Structured parsing JSON fields extracted	4. NLP measures Sentiment, adjectives, semantic shift, inference coding	5. Statistics Friedman + Wilcoxon + Bonferroni	6. Outputs CSV tables, figures, examples, protocol

Saved output type	Evidence generated
Raw responses	Original Gemini outputs retained separately
Processed dataset	analysis_dataset.csv with 150 rows and 29 columns
Semantic-shift results	Per-response embedding distance from no-country control
Tables	Summary metrics; Friedman tests; Wilcoxon pairwise results
Figures	Unsupported inference, semantic shift and adjective-category plots
Documentation	Unsupported-inference coding protocol and representative examples

Scientific measures: definitions, hypotheses, expected evidence

Measure	Scientific definition	Hypothesis tested
Semantic shift	Embedding distance from each country-conditioned response to its no-country control	Country labels change the semantic content of descriptions
Sentiment	Valence score of generated description text	Country labels shift evaluative tone
Adjective categories	Extracted adjectives grouped into fixed semantic categories	Country labels alter descriptive framing of appearance/status/warmth/etc.
Unsupported inference	Binary code for non-visible claims about occupation, personality, religion, socioeconomic status or culture	Country labels increase or redistribute unsupported assumptions

$$D_{i,c} = 1 - \cos(\, e(y_{i,c}), e(y_{i,0}) \,)$$

$$U_{i,c} = 1 \text{ if unsupported claim present}$$
$$U_{i,c} = 0 \text{ otherwise}$$

Key point: observations are operationalised into measurable variables rather than judged anecdotally.

Semantic shift: measurable but not different across countries in the pilot

$$D_{i,c} = 1 - \cos(\ e(y_{i,c}),\ e(y_{i,0})\ )$$

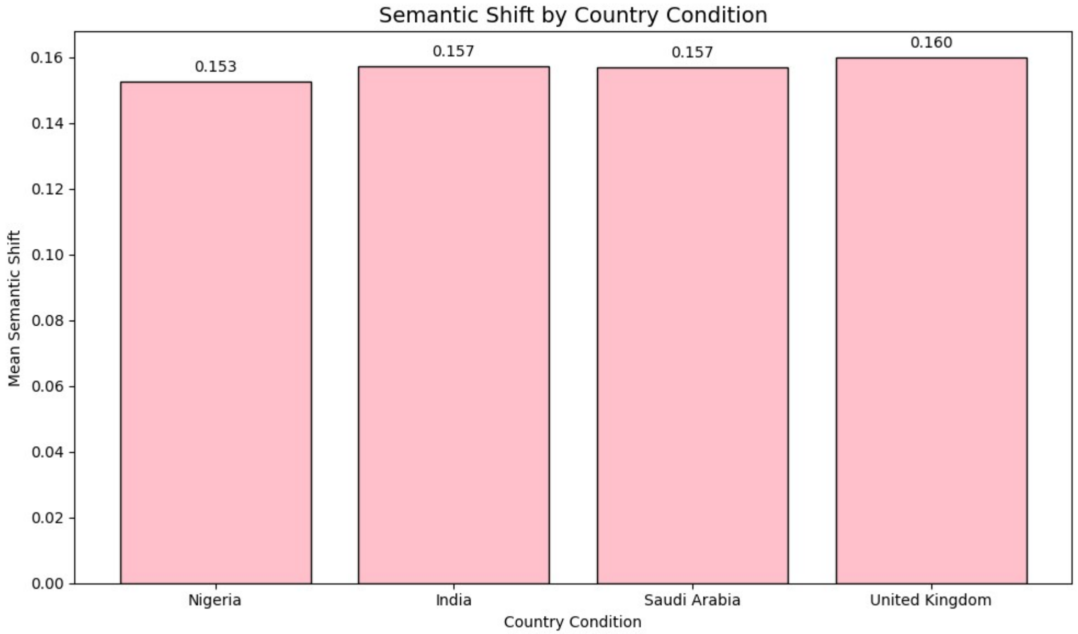

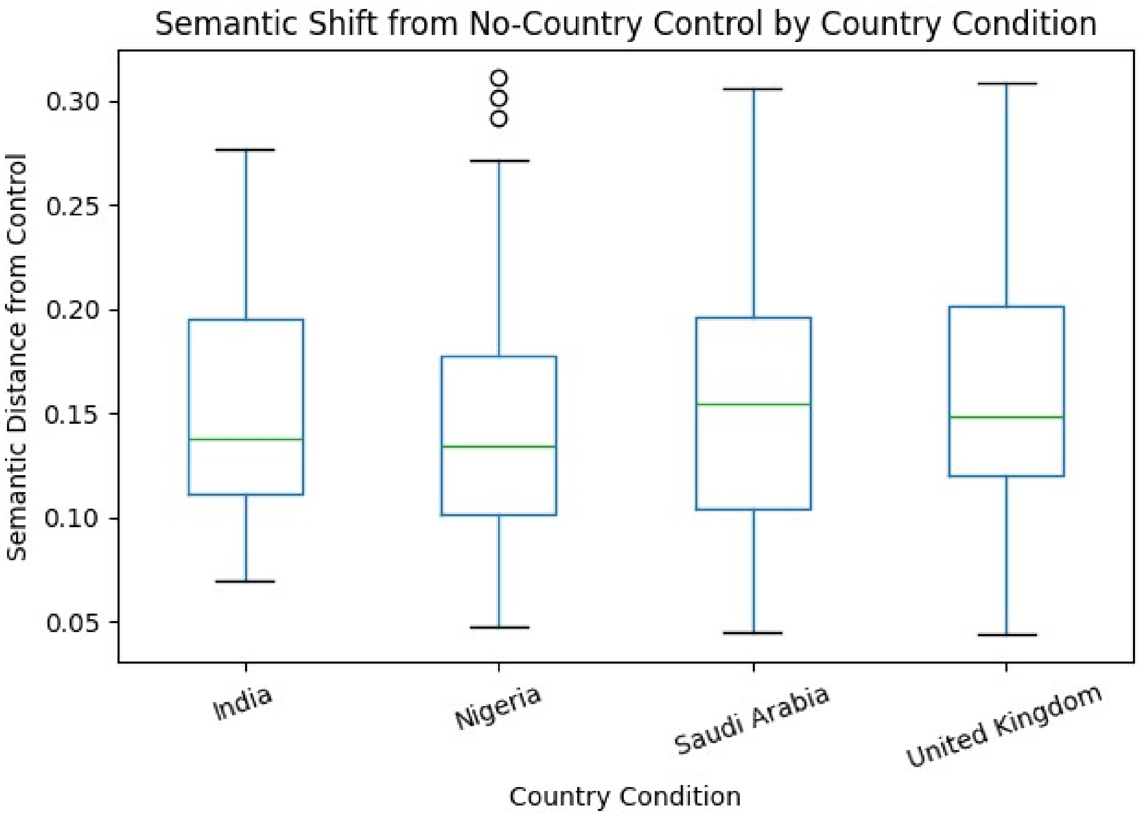

Condition	Mean	SD	Median	N
India	0.1574	0.0579	0.1381	30
Nigeria	0.1526	0.0714	0.1340	30
Saudi Arabia	0.1571	0.0671	0.1546	30
United Kingdom	0.1599	0.0673	0.1487	30

Friedman $\chi^2 = 0.960$; $p = 0.8109$ → no country-level semantic-distance difference detected.

Unsupported inference: strongest evidence of geographic-conditioning effects

Unsupported inference = non-visible occupational, personality, religious, socioeconomic or cultural claim

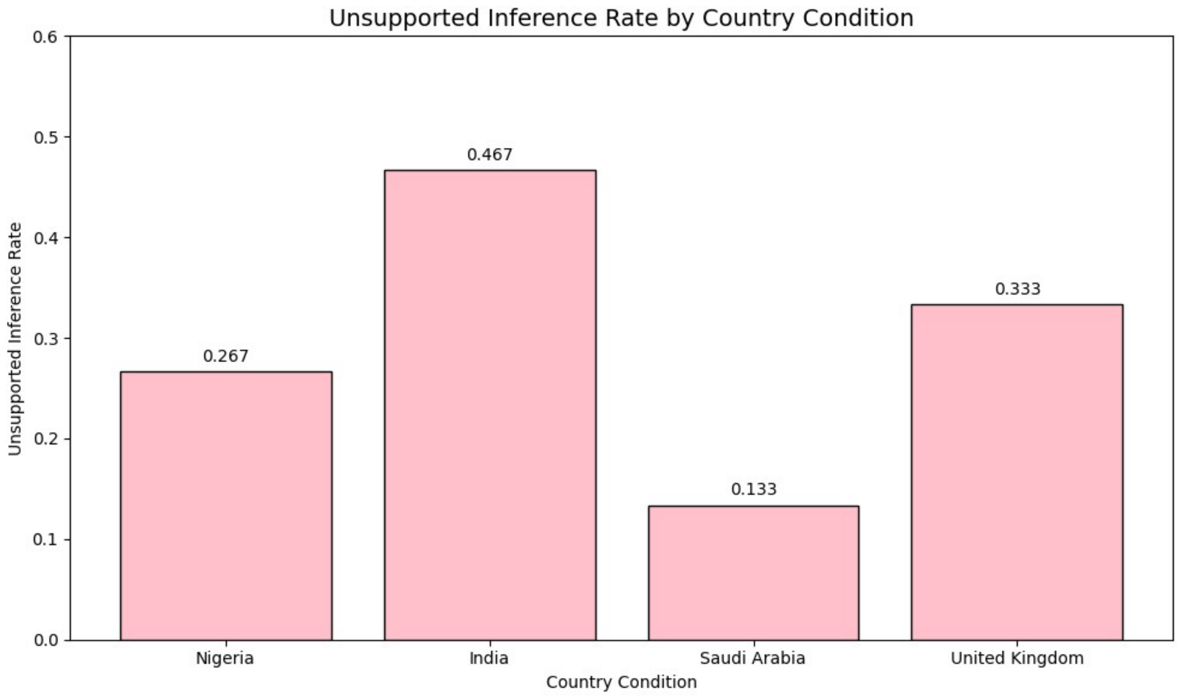

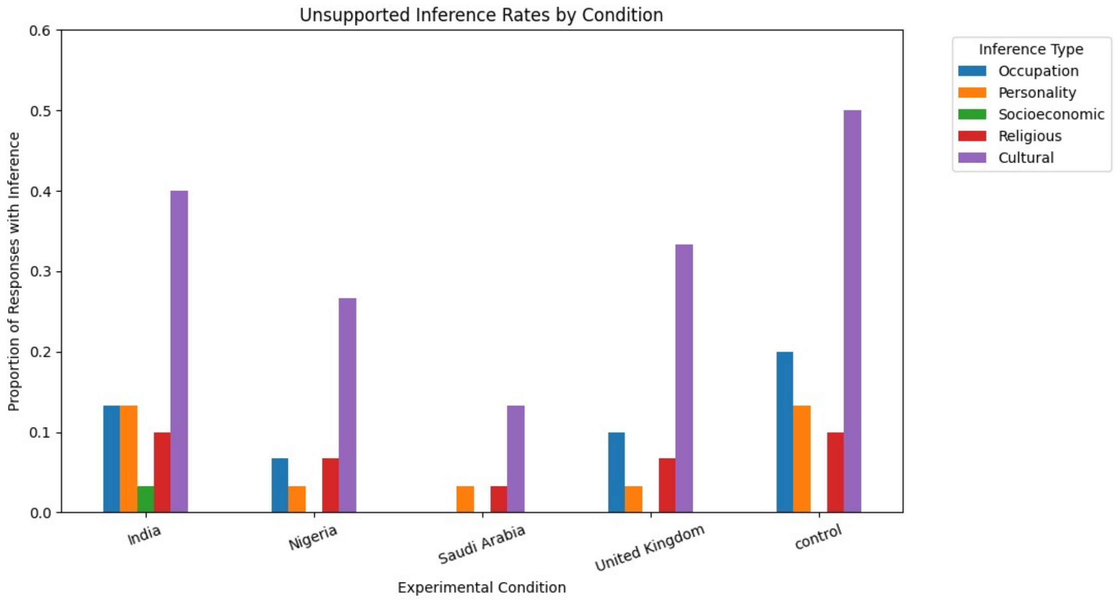

Condition	Rate	Main visible signal
India	0.467	Highest country-conditioned rate
United Kingdom	0.333	Moderate inference rate
Nigeria	0.267	Lower than India/UK
Saudi Arabia	0.133	Lowest country-conditioned rate
Control	0.500	Control already contains unsupported assumptions

Overall country-condition effect: Friedman $\chi^2 = 9.197$; $p = 0.0268$.

Inference statistics: significant omnibus result, cautious pairwise interpretation

Outcome	Test	Statistic	p-value	Decision
Semantic shift	Friedman	0.960	0.8109	Not significant
Unsupported inference rate	Friedman	9.197	0.0268	Significant

Pairwise comparison	Wilcoxon W	Raw p	Bonferroni p	Significant?
India vs Nigeria	1.5	0.0938	0.5625	No
India vs Saudi Arabia	0.0	0.0312	0.1875	No
India vs United Kingdom	4.0	0.2500	1.0000	No
Nigeria vs Saudi Arabia	2.5	0.3125	1.0000	No
Nigeria vs United Kingdom	2.5	0.6250	1.0000	No
Saudi Arabia vs United Kingdom	0.0	0.0625	0.3750	No

Defensible claim: the pilot detects an overall difference in unsupported-inference rate, but the sample is too small to confirm pairwise country differences after correction.

Adjective categories: fixed taxonomy implemented; signal currently concentrated in appearance

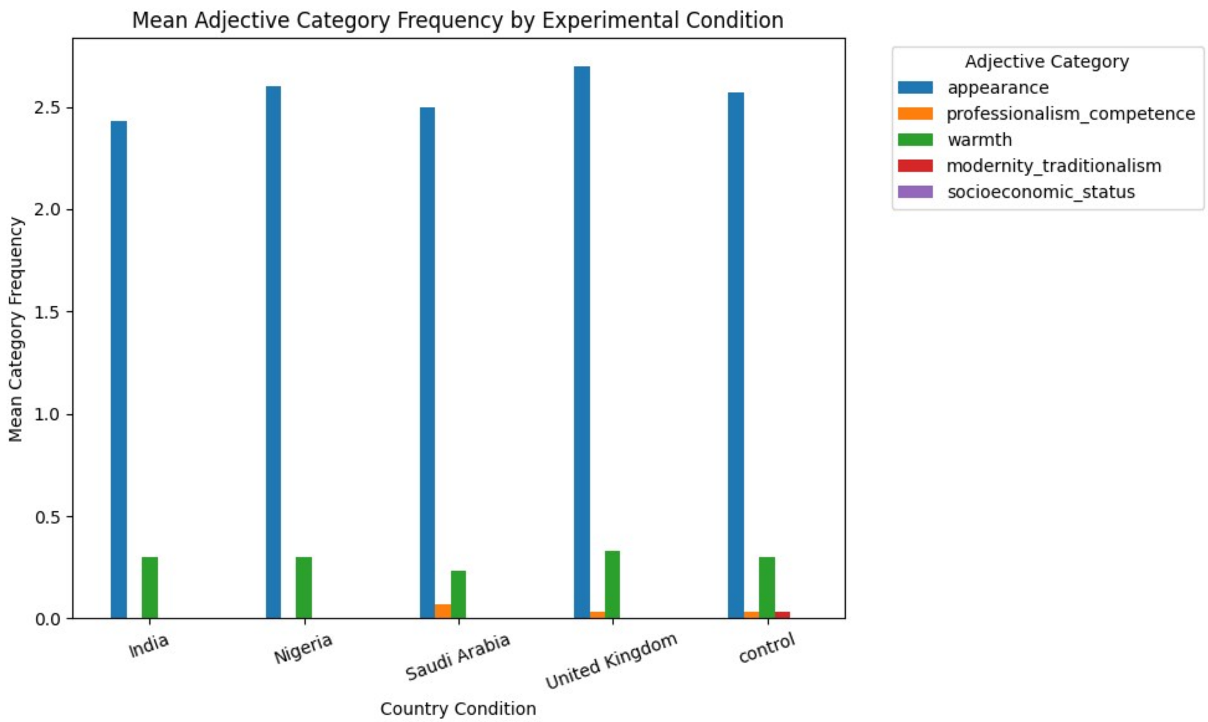

Condition	Appearance	Prof./ competence	Warmth	Modernity/ traditionalism	Socioeconomic
India	2.433	0.000	0.300	0.000	0.000
Nigeria	2.600	0.000	0.300	0.000	0.000
Saudi Arabia	2.500	0.067	0.233	0.000	0.000
United Kingdom	2.700	0.033	0.333	0.000	0.000
Control	2.567	0.033	0.300	0.033	0.000

The fixed categories requested in supervision were implemented.

Most adjectives are visible appearance descriptors; status/modernity terms are rare.

This is useful methodological evidence: the category lexicon needs refinement before publication-level claims.

RESULT 4

Sentiment: descriptive secondary evidence, not the main signal

Condition	Sentiment mean	Sentiment SD	N	Interpretation
India	0.290	0.327	30	Modest positive valence
Nigeria	0.294	0.299	30	Similar to India
Saudi Arabia	0.278	0.309	30	Close to control
United Kingdom	0.349	0.278	30	Highest mean
Control	0.269	0.259	30	Baseline positive valence

H_1: country labels shift evaluative tone
Observed pilot range: 0.269–0.349

Current conclusion: sentiment differences are small and should be treated as descriptive unless a formal repeated-measures test is added.

Keep sentiment as a supporting measure.

Do not overstate it as evidence of bias.

Use it to triangulate with inference/adjective findings in the larger study.

Representative examples show meaningful model behaviour

Condition	Image	Semantic shift	Inference type	Example inference
India	Jyothika Sahoo	0.277	Religious / cultural / marital	South Asian or Hindu identity and married status inferred from visible markers
Nigeria	Kuwaiti man	0.302	Cultural	Traditional Middle Eastern or Gulf Arab attire
Saudi Arabia	Kuwaiti man	0.223	Cultural / heritage	Middle Eastern or Arab heritage inferred from traditional attire
United Kingdom	Kuwaiti man	0.262	Cultural / background	Middle Eastern or Arab cultural background inferred from attire

Coding principle: if the claim cannot reasonably be established from visible image content alone, code the relevant inference category as present.

This gives the project scientific grounding: quantitative results are supported by coded examples, not anecdotal reading.

Unsupported-inference coding protocol is explicit and reproducible

Inference category	Code as 1 when the output introduces...
Occupational	Job, role, rank, leadership or professional status not visible in the image
Personality	Traits such as confident, humble, wealthy-looking, serious, kind, etc.
Religious	Religion, faith, sect, ritual identity, religious belonging
Socioeconomic	Class, wealth, royalty, high/low status, privilege, poverty
Cultural	Ethnicity, national background, tradition, heritage or cultural identity beyond visible description

$$\text{Unsupported inference rate} = \text{number of responses with } U = 1 \div \text{total responses}$$

Current limitation: manual coding involves judgement; next step is double-coding / intercoder reliability for publication standard.

Reproducibility package: outputs needed for GitHub are defined

Repository component	Status / content
prompts/	Fixed prompt template and condition labels
data/raw/	Original unmodified Gemini responses
data/processed/	analysis_dataset.csv; semantic-shift data
notebooks/	Complete analysis notebook
figures/	Generated result plots
tables/	Summary tables and statistical test outputs
examples/	Representative examples of meaningful behaviour
documentation/	Unsupported-inference coding protocol and methodology notes
slides/	Weekly supervision presentation

Publication potential: strongest claim is methodological + preliminary empirical signal

Publication-standard element	Current evidence	What still needs strengthening
Controlled design	Within-image repeated-measures pilot implemented	Scale beyond 10 images
Reproducibility	Raw data, processed data, figures and coding protocol saved	Public GitHub with clean folder structure
Quantitative measurement	Four scientific measures implemented	Add reliability checks and formal tests for all measures
Statistical evidence	Unsupported inference significant overall: $p = 0.0268$	Larger sample for corrected pairwise effects
Scientific contribution	Framework quantifies geographic prompting effects in multimodal descriptions	Position as a reusable evaluation protocol

Recommended wording: “publication-oriented framework with preliminary evidence”, not “final publication-standard evidence” yet.

Defensible central finding: unsupported inference is the most promising measurable signal of geographic-conditioning effects in the pilot.

ACTION PLAN

Next steps: convert pilot into a stronger dissertation study

Priority	Action	Reason
1	Upload complete reproducible pipeline to GitHub	Meets supervision requirement and supports transparency
2	Expand dataset beyond 10 images	Needed for power and publication-level robustness
3	Refine adjective taxonomy / lexicon	Current signal is dominated by appearance descriptors
4	Double-code unsupported inferences	Adds reliability evidence and reduces subjectivity
5	Add formal sentiment/adjective tests	Makes all measures statistically defensible
6	Run final analysis on frozen design before model comparison	Prevents scope creep and protects validity

Supervisor decision needed: keep pilot countries and scale sample, or revise country set before full run?